

Language-Table ablation — validation action-classification accuracy
(3,000 simulated episodes · 90/10 split · 3 seeds each · random baseline = 12.5%)

